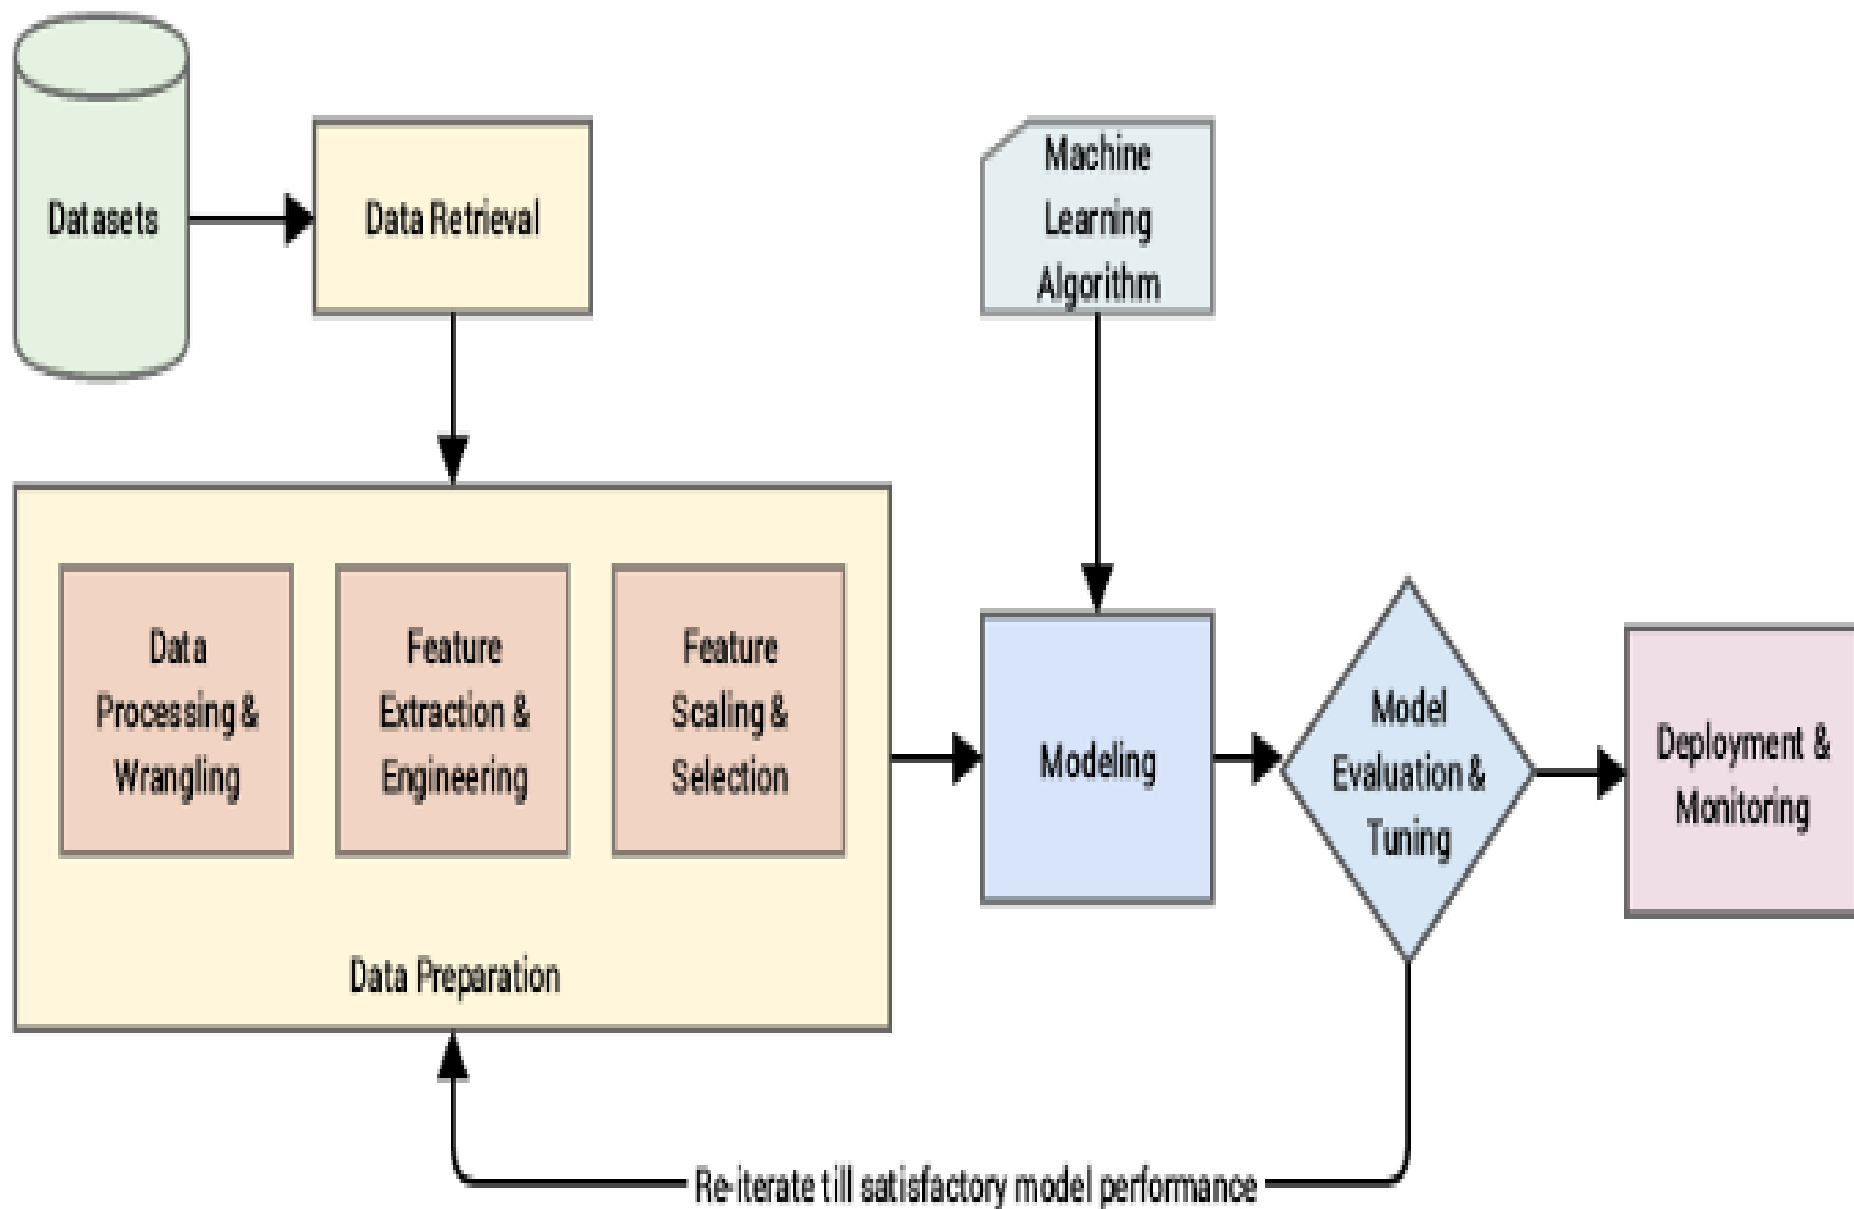



Feature Engineering for Machine Learning



A standard machine learning pipeline (source: Practical Machine Learning with Python, Apress/Springer)

What data scientists spend the most time doing



- Building training sets: 3%
- Cleaning and organizing data: 60%
- Collecting data sets: 19%
- Mining data for patterns: 9%
- Refining algorithms: 4%
- Other: 5%

Source: <https://www.forbes.com/sites/gilpress/2016/03/23/data-preparation-most-time-consuming-least-enjoyable-data-science-task-survey-says/>

Motivation

“Coming up with features is difficult, time-consuming, requires expert knowledge. ‘Applied machine learning’ is basically feature engineering.”

— *Prof. Andrew Ng.*

*“Feature engineering is the process of transforming **raw data** into **features** that better represent **the underlying problem** to **the predictive models**, resulting in improved **model accuracy** on **unseen data**.”*

— *Dr. Jason Brownlee*

“At the end of the day, some machine learning projects succeed and some fail. What makes the difference? Easily the most important factor is the features used.”

— *Prof. Pedro Domingos*

Feature engineering is another topic which doesn't seem to merit any review papers or books, or even chapters in books, but it is absolutely vital to ML success. [...]

Much of the success of machine learning is actually success in engineering features that a learner can understand.

— Scott Locklin, in “Neglected machine learning ideas”

WHY ?

- 1. Make simple model to perform much better than complex model.**
- 2. Reduce model selection time.**
- 3. Reduce training time by simplifying the model.**

PANDAS

- Pandas is an open-source, BSD-licensed Python library providing high-performance, easy-to-use data structures and data analysis tools for the Python programming language.

Key Features of Pandas

- Fast and efficient DataFrame object with default and customized indexing.
- Tools for loading data into in-memory data objects from different file formats.
- Data alignment and integrated handling of missing data.
- Reshaping and pivoting of date sets.
- Label-based slicing, indexing and subsetting of large data sets.
- Columns from a data structure can be deleted or inserted.
- Group by data for aggregation and transformations.
- High performance merging and joining of data.
- Time Series functionality.

Pandas Data Structures

- Series

Series is a one-dimensional array like structure with homogeneous data. For example, the following series is a collection of integers 10, 23, 56,

- Key Points
- Homogeneous data
- Size Immutable
- Values of Data Mutable

Pandas Data Structures

DataFrame

DataFrame is a two-dimensional array with heterogeneous data

Example

Name	Age	Gender	Rating
Steve	32	Male	3.45
Lia	28	Female	4.6
Vin	45	Male	3.9
Katie	38	Female	2.78

Pandas Data Structures

Key Points

- Heterogeneous data
- Size Mutable
- Data Mutable

Panel

Panel

- Panel is a three-dimensional data structure with heterogeneous data. It is hard to represent the panel in graphical representation. But a panel can be illustrated as a container of DataFrame.
- Key Points
- Heterogeneous data
- Size Mutable
- Data Mutable

Data Selection

- **Data Import**

Use these commands to import data from a variety of different sources and formats.

- `pd.read_csv(filename)` | From a CSV file
- `pd.read_table(filename)` | From a delimited text file (like TSV)
- `pd.read_excel(filename)` | From an Excel file
- `pd.read_sql(query, connection_object)` | Read from a SQL table/database
- `pd.read_json(json_string)` | Read from a JSON formatted string, URL or file.
- `pd.read_html(url)` | Parses an html URL, string or file and extracts tables to a list of dataframes
- `pd.read_clipboard()` | Takes the contents of your clipboard and passes it to `read_table()`
- `pd.DataFrame(dict)` | From a dict, keys for columns names, values for data as lists

Data Export

Use these commands to export a DataFrame to CSV, .xlsx, SQL, or JSON.

- `df.to_csv(filename)` | Write to a CSV file
- `df.to_excel(filename)` | Write to an Excel file
- `df.to_sql(table_name, connection_object)` | Write to a SQL table
- `df.to_json(filename)` | Write to a file in JSON format

Viewing/Inspecting Data

Use these commands to take a look at specific sections of your pandas DataFrame or Series.

- `df.head(n)` | First `n` rows of the DataFrame
- `df.tail(n)` | Last `n` rows of the DataFrame
- `df.shape` | Number of rows and columns
- `df.info()` | Index, Datatype and Memory information
- `df.describe()` | Summary statistics for numerical columns
- `s.value_counts(dropna=False)` | View unique values and counts
- `df.apply(pd.Series.value_counts)` | Unique values and counts for all columns

Selection

Use these commands to select a specific subset of your data.

- `df[col]` | Returns column with label `col` as Series
- `df[[col1, col2]]` | Returns columns as a new DataFrame
- `s.iloc[0]` | Selection by position
- `s.loc['index_one']` | Selection by index
- `df.iloc[0,:]` | First row
- `df.iloc[0,0]` | First element of first column

Data Cleaning

- Use these commands to perform a variety of data cleaning tasks.
- `df.columns = ['a','b','c']` | Rename columns
- `pd.isnull()` | Checks for null Values, Returns Boolean Array
- `pd.notnull()` | Opposite of `pd.isnull()`
- `df.dropna()` | Drop all rows that contain null values
- `df.dropna(axis=1)` | Drop all columns that contain null values
- `df.dropna(axis=1,thresh=n)` | Drop all rows have have less than n non null values
- `df.fillna(x)` | Replace all null values with x
- `s.fillna(s.mean())` | Replace all null values with the mean (mean can be replaced with almost any function from the statistics module)
- `s.astype(float)` | Convert the datatype of the series to float
- `s.replace(1,'one')` | Replace all values equal to 1 with 'one'
- `s.replace([1,3],['one','three'])` | Replace all 1 with 'one' and 3 with 'three'
- `df.rename(columns=lambda x: x + 1)` | Mass renaming of columns
- `df.rename(columns={'old_name': 'new_name'})` | Selective renaming
- `df.set_index('column_one')` | Change the index
- `df.rename(index=lambda x: x + 1)` | Mass renaming of index

Filter, Sort, and Groupby

Use these commands to filter, sort, and group your data.

- `df[df[col] > 0.5]` | Rows where the column col is greater than 0.5
- `df[(df[col] > 0.5) & (df[col] < 0.7)]` | Rows where $0.7 > \text{col} > 0.5$
- `df.sort_values(col1)` | Sort values by col1 in ascending order
- `df.sort_values(col2, ascending=False)` | Sort values by col2 in descending order
- `df.sort_values([col1, col2], ascending=[True, False])` | Sort values by col1 in ascending order then col2 in descending order
- `df.groupby(col)` | Returns a groupby object for values from one column
- `df.groupby([col1, col2])` | Returns groupby object for values from multiple columns
- `df.groupby(col1)[col2]` | Returns the mean of the values in col2, grouped by the values in col1 (mean can be replaced with almost any function from the statistics module)
- `df.pivot_table(index=col1, values=[col2, col3], aggfunc=mean)` | Create a pivot table that groups by col1 and calculates the mean of col2 and col3
- `df.groupby(col1).agg(np.mean)` | Find the average across all columns for every unique col1 group
- `df.apply(np.mean)` | Apply the function `np.mean()` across each column
- `df.apply(np.max, axis=1)` | Apply the function `np.max()` across each row

Statistics

Use these commands to perform various statistical tests. (These can all be applied to a series as well.)

- `df.describe()` | Summary statistics for numerical columns
- `df.mean()` | Returns the mean of all columns
- `df.corr()` | Returns the correlation between columns in a DataFrame
- `df.count()` | Returns the number of non-null values in each DataFrame column
- `df.max()` | Returns the highest value in each column
- `df.min()` | Returns the lowest value in each column
- `df.median()` | Returns the median of each column
- `df.std()` | Returns the standard deviation of each column

Feature Engineering

- Managing missing features
- Managing Categorical Features
- Data Scaling and Normalization
- Whitening
- Dimensionality reduction
- Creating training and testing data sets

CASE STUDY

- Print first and last five rows
- Clean data and update the CSV file
- Find the most expensive car company name
- Print All Toyota Cars details
- Sort all cars by Price column

Categorical data

- Binary: A variable that has only 2 values. For example, True/False or Yes/No.
- Ordinal: A variable that has some order associated with it like our place example above.
- Nominal: A variable that has no numerical importance, for example color or city.

Managing Categorical Data

- Encoding or continuation is the transformation of categorical variables to binary or numerical counterparts.
- Types
 - Binary
 - Target-based-ordinal, one-hot encoding

Data Scaling and Normalization

